

High Performance Computing

COM403P

Week-5

Vector Addition

Kadali Paul Babu

CED19I002

Objective

Vector Addition for given $n \times n$ double precision floating point numbers.

Serial Code

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#include <time.h>

#define SIZE 100000

int main()
{
    double a[SIZE], b[SIZE], c[SIZE], rand_a, rand_b;
    double start, end, exec;

    start = omp_get_wtime();

    for (int i = 0; i < SIZE; i++)
    {
        rand_a = rand();
        rand_b = rand();

        a[i] = i*rand_a;
        b[i] = i*rand_b;

        c[i] = a[i] + b[i];
    }

    end = omp_get_wtime();

    exec = end - start;

    printf("Serial Exec time - %f\n", exec);

    return 0;
}
```

Parallel Code

```
#include <stdio.h>
#include <stdlib.h>
#include <mpi.h>
#include <time.h>

#define MASTER 0
#define FROM_MASTER 1
#define FROM_WORKER 2

#define SIZE 100

int main()
{
    MPI_Init(NULL, NULL);

    long double a[SIZE][SIZE], b[SIZE][SIZE], c[SIZE][SIZE], rand_a, rand_b;
    double start, end, exec;
    int i,j,k;

    int avgrow, extra;
    int offset, mtype;
    int dest, rows;
    int source;
    MPI_Status status;

    start = MPI_Wtime();

    int taskid;
    MPI_Comm_rank(MPI_COMM_WORLD,&taskid);

    int numtasks;
    MPI_Comm_size(MPI_COMM_WORLD,&numtasks);

    int workers = numtasks -1;

    if(taskid == MASTER)
    {
        for (i = 0; i < SIZE; i++)
        {
            a[i][j] = (i+j)*1.22;
            b[i][j] = (i+j)*1.22;
        }
    }
}
```

```
//-----  
-----  
    // Master sending Data to Worker tasks  
  
    avgrow = SIZE/workers;  
    extra = SIZE%workers;  
  
    offset = 0;  
    mtype = FROM_MASTER;  
  
    for(dest = 1; dest <= workers; dest++)  
    {  
        rows = (dest <= extra)?avgrow+1:avgrow;  
  
        // printf("Sending %d rows to task-%d offset=%d\n",rows,dest,offset);  
  
        MPI_Send(&offset, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);  
        MPI_Send(&rows, 1, MPI_INT, dest, mtype, MPI_COMM_WORLD);  
        MPI_Send(&a[offset], rows, MPI_LONG_DOUBLE, dest, mtype, MPI_COMM_WORLD);  
        MPI_Send(&b[offset], rows, MPI_LONG_DOUBLE, dest, mtype, MPI_COMM_WORLD);  
  
        offset += rows;  
    }  
  
//-----  
-----  
    //Receiving results from Workers  
  
    mtype = FROM_WORKER;  
    for(i = 1; i <= workers; i++)  
    {  
        source = i;  
  
        MPI_Recv(&offset, 1, MPI_INT, source, mtype, MPI_COMM_WORLD, &status);  
        MPI_Recv(&rows, 1, MPI_INT, source, mtype, MPI_COMM_WORLD, &status);  
        MPI_Recv(&c[offset], rows, MPI_LONG_DOUBLE, source, mtype, MPI_COMM_WORLD,  
&status);  
  
        // printf("Received results from task=%d\n", source);  
    }  
    // printf("Result Matrix\n");
```

```
// for(i = 0; i < SIZE;i++)
// {
//     // printf("\n");
//     for(int j = 0; j < SIZE;j++)
//     {
//         // printf("%6.2Lf  ",c[i][j]);
//     }
// }

end = MPI_Wtime();

exec = end - start;
printf("MPI Exec time - %f\n", exec);
}
```

```
//-----
-----
```

```
//-----
-----
```

```
//Worker Task
if(taskid > MASTER)
{
    mtype = FROM_MASTER;

    MPI_Recv(&offset, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,&status);
    MPI_Recv(&rows, 1, MPI_INT, MASTER, mtype, MPI_COMM_WORLD,&status);
    MPI_Recv(&a, rows, MPI_LONG_DOUBLE, MASTER, mtype, MPI_COMM_WORLD,&status);
    MPI_Recv(&b, rows, MPI_LONG_DOUBLE, MASTER, mtype, MPI_COMM_WORLD,&status);

    for(i = 0; i < rows; i++)
    {
        c[i] = a[i] + b[i];
    }

    mtype = FROM_WORKER;


    MPI_Send(&offset,1, MPI_INT, MASTER,mtype, MPI_COMM_WORLD);
    MPI_Send(&rows,1, MPI_INT, MASTER,mtype, MPI_COMM_WORLD);
}
```

```
MPI_Send(&c, rows, MPI_LONG_DOUBLE, dest, mtype, MPI_COMM_WORLD);  
  
}  
MPI_Finalize();  
  
return 0;  
}
```

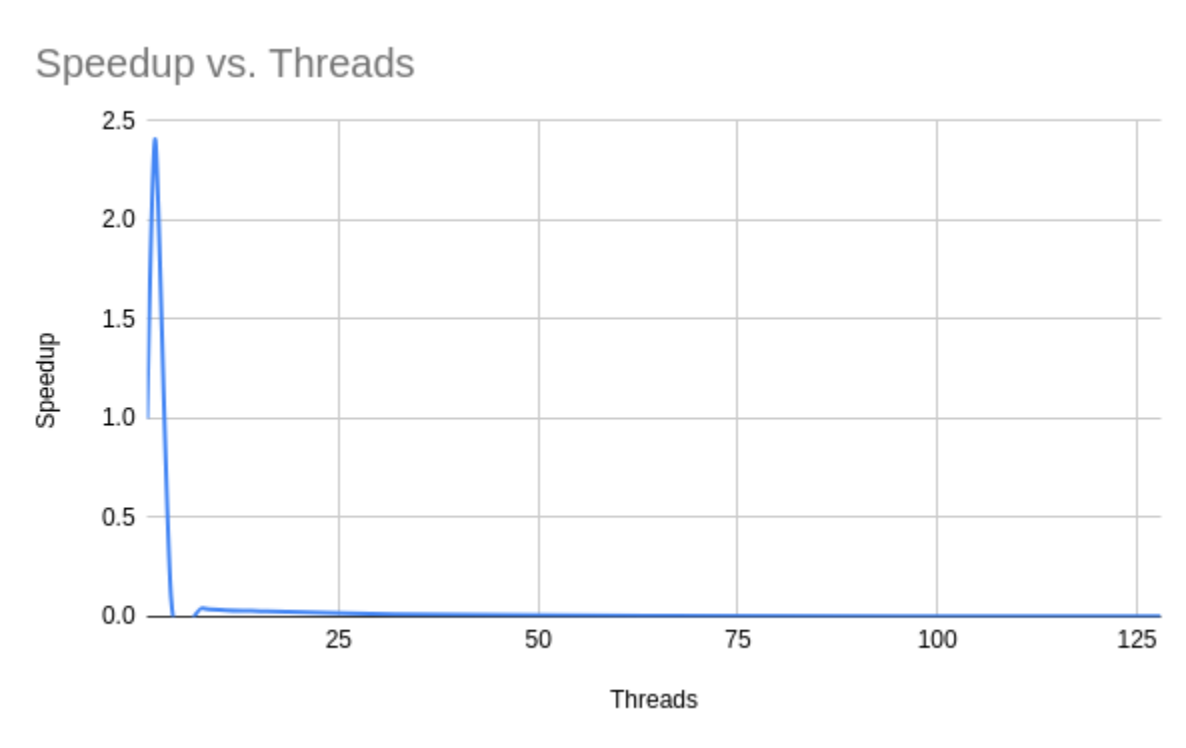
Observations

| no of processors | Exec Time | Speedup | Parallelization Factor |
|------------------|-----------|----------------|------------------------|
| 1 | 0.001586 | 1 | 0 |
| 2 | 0.000659 | 2.406676783 | 0.487074401 |
| 4 | 0.016021 | 0.09899506897 | -6.067675494 |
| 8 | 0.036504 | 0.04344729345 | -7.338797814 |
| 16 | 0.060079 | 0.02639857521 | 12.29361076 |
| 32 | 0.128724 | 0.01232093471 | 133.6044557 |
| 64 | 0.317171 | 0.005000457167 | 862.2540984 |
| 128 | 1.252296 | 0.001266473741 | 7623.074821 |

Threads vs Time



Speedups vs Threads



Inferences

Since MPI is a distributed memory architecture, the communication overhead between nodes causes the parallel code to run slower compared to serial code (running in 1 node or only in master)